

Importing Libraries

```
In [ ]:
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

Reading the data

```
In [ ]:
smoke_data=pd.read_csv("/content/drive/MyDrive/CSV files/train_dataset.csv")
```

Analysing the Data

```
In [ ]:
smoke_data.head()
```

Out[]:

	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	systolic	relaxation	...	HD
0	35	170	85	97.0	0.9	0.9	1	1	118	78	...	7
1	20	175	110	110.0	0.7	0.9	1	1	119	79	...	7
2	45	155	65	86.0	0.9	0.9	1	1	110	80	...	5
3	45	165	80	94.0	0.8	0.7	1	1	158	88	...	4
4	20	165	60	81.0	1.5	0.1	1	1	109	64	...	4

5 rows x 23 columns



```
In [ ]:
smoke_data.tail()
```

Out[]:

	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	systolic	relaxation	...
38979	40	165	60	80.0	0.4	0.6	1	1	107	60	...
38980	45	155	55	75.0	1.5	1.2	1	1	126	72	...
38981	40	170	105	124.0	0.6	0.5	1	1	141	85	...
38982	40	160	55	75.0	1.5	1.5	1	1	95	69	...
38983	55	175	60	81.1	1.0	1.0	1	1	114	66	...

5 rows x 23 columns



```
In [ ]:
smoke_data.shape
```

Out[]:

(38984, 23)

In []:

```
smoke_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38984 entries, 0 to 38983
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   38984 non-null  int64
1   height(cm)                           38984 non-null  int64
2   weight(kg)                            38984 non-null  int64
3   waist(cm)                             38984 non-null  float64
4   eyesight(left)                        38984 non-null  float64
5   eyesight(right)                       38984 non-null  float64
6   hearing(left)                         38984 non-null  int64
7   hearing(right)                        38984 non-null  int64
8   systolic                              38984 non-null  int64
9   relaxation                            38984 non-null  int64
10  fasting blood sugar                   38984 non-null  int64
11  Cholesterol                           38984 non-null  int64
12  triglyceride                          38984 non-null  int64
13  HDL                                   38984 non-null  int64
14  LDL                                   38984 non-null  int64
15  hemoglobin                            38984 non-null  float64
16  Urine protein                         38984 non-null  int64
17  serum creatinine                      38984 non-null  float64
18  AST                                   38984 non-null  int64
19  ALT                                   38984 non-null  int64
20  Gtp                                    38984 non-null  int64
21  dental caries                         38984 non-null  int64
22  smoking                               38984 non-null  int64
dtypes: float64(5), int64(18)
memory usage: 6.8 MB
```

In []:

```
smoke_data.isnull().sum()
```

Out[]:

```
age                0
height(cm)         0
weight(kg)          0
waist(cm)           0
eyesight(left)      0
eyesight(right)     0
hearing(left)       0
hearing(right)      0
systolic            0
relaxation          0
fasting blood sugar 0
Cholesterol         0
triglyceride        0
HDL                 0
LDL                 0
hemoglobin          0
Urine protein       0
serum creatinine    0
AST                 0
ALT                 0
Gtp                 0
dental caries       0
smoking             0
dtype: int64
```

In []:

```
smoke_data.describe()
```

Out[]:

	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	
count	38984.000000	38984.000000	38984.000000	38984.000000	38984.000000	38984.000000	38984.000000	38984.000000	38984.000000
mean	44.127591	164.689488	65.938718	82.062115	1.014955	1.008768	1.025369	1.026190	
std	12.063564	9.187507	12.896581	9.326798	0.498527	0.493813	0.157246	0.159703	
min	20.000000	130.000000	30.000000	51.000000	0.100000	0.100000	1.000000	1.000000	
25%	40.000000	160.000000	55.000000	76.000000	0.800000	0.800000	1.000000	1.000000	
50%	40.000000	165.000000	65.000000	82.000000	1.000000	1.000000	1.000000	1.000000	
75%	55.000000	170.000000	75.000000	88.000000	1.200000	1.200000	1.000000	1.000000	
max	85.000000	190.000000	135.000000	129.000000	9.900000	9.900000	2.000000	2.000000	

8 rows x 23 columns

In []:

```
smoke_data['smoking'].value_counts()
```

Out[]:

0 24666
1 14318
Name: smoking, dtype: int64

In []:

```
X= smoke_data.drop(columns='smoking',axis=1)  
Y= smoke_data['smoking']
```

In []:

```
print(X)
```

	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	systolic	relaxation	triglyceride	HDL	LDL	hemoglobin	Urine protein
0	35	170	85	97.0	0.9	0.9	1	1	118	78	153	70	142	19.8	1
1	20	175	110	110.0	0.7	0.9	1	1	119	79	128	71	114	15.9	1
2	45	155	65	86.0	0.9	0.9	1	1	110	80	120	57	112	13.7	3
3	45	165	80	94.0	0.8	0.7	1	1	158	88	366	46	91	16.9	1
4	20	165	60	81.0	1.5	0.1	1	1	109	64	200	47	92	14.9	1
...
38979	40	165	60	80.0	0.4	0.6	1	1	107	60
38980	45	155	55	75.0	1.5	1.2	1	1	126	72
38981	40	170	105	124.0	0.6	0.5	1	1	141	85
38982	40	160	55	75.0	1.5	1.5	1	1	95	69
38983	55	175	60	81.1	1.0	1.0	1	1	114	66

38979	...	53	61	72	12.3	1
38980	...	100	76	131	12.5	2
38981	...	196	48	138	17.1	1
38982	...	48	79	116	12.0	1
38983	...	57	64	137	13.9	1

	serum creatinine	AST	ALT	Gtp	dental caries
0	1.0	61	115	125	1
1	1.1	19	25	30	1
2	0.6	1090	1400	276	0
3	0.9	32	36	36	0
4	1.2	26	28	15	0
...
38979	0.5	18	18	21	1
38980	0.6	23	11	12	0
38981	0.8	24	23	35	1
38982	0.6	24	20	17	0
38983	1.0	18	12	16	0

[38984 rows x 22 columns]

In []:

```
print(Y)
```

```
0      1
1      0
2      0
3      0
4      0
...
38979  0
38980  0
38981  1
38982  1
38983  1
```

Name: smoking, Length: 38984, dtype: int64

Training the Data

In []:

```
X_train, X_test, Y_train, Y_test= train_test_split(X, Y, test_size=0.4, stratify=Y, random_s
tate=4)
```

In []:

```
print(X.shape, X_train.shape, X_test.shape)
```

```
(38984, 22) (23390, 22) (15594, 22)
```

Logistic Regression

In []:

```
model = LogisticRegression()
```

In []:

```
model.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: Convergence
Warning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
```

Out []:

```
LogisticRegression()
```

In []:

```
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

In []:

```
print('Accuracy on Training data : ', training_data_accuracy)
```

Accuracy on Training data : 0.7114151346729372

In []:

```
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

In []:

```
print('Accuracy on Test data : ', test_data_accuracy)
```

Accuracy on Test data : 0.7123252533025523

In []:

```
input_data = (30,180,90,94.0,1.0,0.8,1,1,115,72,88,177,103,53,103,13.5,1,1.0,19,29,30,0)
```

```
# change the input data to a numpy array
```

```
input_data_as_numpy_array= np.asarray(input_data)
```

```
# reshape the numpy array as we are predicting for only on instance
```

```
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
```

```
prediction = model.predict(input_data_reshaped)
```

```
print(prediction)
```

```
if (prediction[0]== 0):
```

```
    print('The Person does not Smoke')
```

```
else:
```

```
    print('The Person Smokes')
```

```
[0]
```

The Person does not Smoke

/usr/local/lib/python3.7/dist-packages/sklearn/base.py:451: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
"X does not have valid feature names, but"

Implementing Test Data on the created model

In []:

```
new_test_data = pd.read_csv("/content/drive/MyDrive/CSV files/test_dataset.csv")
new_test_data.head()
```

Out []:

	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	systolic	relaxation	...	trig
0	40	170	65	75.1	1.0	0.9	1	1	120	70	...	
1	45	170	75	89.0	0.7	1.2	1	1	100	67	...	
2	30	180	90	94.0	1.0	0.8	1	1	115	72	...	
3	60	170	50	73.0	0.5	0.7	1	1	118	78	...	
4	30	170	65	78.0	1.5	1.0	1	1	110	70	...	

5 rows x 22 columns

In []:

```
new_test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16708 entries, 0 to 16707
Data columns (total 22 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   16708 non-null  int64
 1   height(cm)           16708 non-null  int64
 2   weight(kg)           16708 non-null  int64
 3   waist(cm)            16708 non-null  float64
 4   eyesight(left)       16708 non-null  float64
 5   eyesight(right)      16708 non-null  float64
 6   hearing(left)        16708 non-null  int64
 7   hearing(right)       16708 non-null  int64
 8   systolic             16708 non-null  int64
 9   relaxation            16708 non-null  int64
10   fasting blood sugar  16708 non-null  int64
11   Cholesterol          16708 non-null  int64
12   triglyceride         16708 non-null  int64
13   HDL                  16708 non-null  int64
14   LDL                  16708 non-null  int64
15   hemoglobin           16708 non-null  float64
16   Urine protein        16708 non-null  int64
17   serum creatinine     16708 non-null  float64
18   AST                  16708 non-null  int64
19   ALT                  16708 non-null  int64
20   Gtp                  16708 non-null  int64
21   dental caries        16708 non-null  int64
dtypes: float64(5), int64(17)
memory usage: 2.8 MB
```

In []:

```
input_data=new_test_data
prediction = model.predict(input_data)
#print("\n",prediction)
res=pd.DataFrame(prediction)
res.index=new_test_data.index
res.columns=["prediction"]
res.to_csv("prediction_results_smoke_data.csv")
#files.download("prediction_results_smoke_data.csv")
#if (prediction[0]== 0):
# print('The Person does not Smoke')
#else:
# print('The Person Smokes')
```

Downloading the Prediction file

In []:

```
from google.colab import files
files.download("prediction_results_smoke_data.csv")
```

Random Forest Classifier

In [88]:

```
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import accuracy_score, f1_score
from sklearn.ensemble import RandomForestClassifier
```

```
In [95]:
```

```
forest = RandomForestClassifier(random_state=1, n_estimators=1000, max_depth=5)

forest.fit(X_train, Y_train)
```

```
Out[95]:
```

```
RandomForestClassifier(max_depth=5, n_estimators=1000, random_state=1)
```

Boruta

```
In [ ]:
```

```
!pip install boruta
```

```
Collecting boruta
```

```
  Downloading Boruta-0.3-py3-none-any.whl (56 kB)
```

```
    |████████████████████| 56 kB 2.6 MB/s
```

```
Requirement already satisfied: numpy>=1.10.4 in /usr/local/lib/python3.7/dist-packages (from boruta) (1.21.6)
```

```
Requirement already satisfied: scikit-learn>=0.17.1 in /usr/local/lib/python3.7/dist-packages (from boruta) (1.0.2)
```

```
Requirement already satisfied: scipy>=0.17.0 in /usr/local/lib/python3.7/dist-packages (from boruta) (1.4.1)
```

```
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.17.1->boruta) (1.1.0)
```

```
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.17.1->boruta) (3.1.0)
```

```
Installing collected packages: boruta
```

```
Successfully installed boruta-0.3
```

```
In [104]:
```

```
from boruta import BorutaPy
```

```
In [111]:
```

```
boruta_selector = BorutaPy(forest, n_estimators='auto', verbose=2, random_state=1)  # initialize the boruta selector
boruta_selector.fit(np.array(X_train), np.array(Y_train))  #
```

```
Iteration: 1 / 100
```

```
Confirmed: 0
```

```
Tentative: 22
```

```
Rejected: 0
```

```
Iteration: 2 / 100
```

```
Confirmed: 0
```

```
Tentative: 22
```

```
Rejected: 0
```

```
Iteration: 3 / 100
```

```
Confirmed: 0
```

```
Tentative: 22
```

```
Rejected: 0
```

```
Iteration: 4 / 100
```

```
Confirmed: 0
```

```
Tentative: 22
```

```
Rejected: 0
```

```
Iteration: 5 / 100
```

```
Confirmed: 0
```

```
Tentative: 22
```

```
Rejected: 0
```

```
Iteration: 6 / 100
```

```
Confirmed: 0
```

```
Tentative: 22
```

```
Rejected: 0
```

```
Iteration: 7 / 100
```

```
Confirmed: 0
```

```
Tentative: 22
```

```
Rejected: 0
```

```
Iteration: 8 / 100
```

```
Confirmed: 19
```

```
Tentative: 0
```

Tentative: 0
Rejected: 3

BorutaPy finished running.

Iteration: 9 / 100
Confirmed: 19
Tentative: 0
Rejected: 3

Out[111]:

```
BorutaPy(estimator=RandomForestClassifier(max_depth=5, n_estimators=132,  
                                           random_state=RandomState(MT19937) at 0x7FC0774  
E1380),  
         n_estimators='auto',  
         random_state=RandomState(MT19937) at 0x7FC0774E1380, verbose=2)
```

In [112]:

```
print("Selected Features: ", boruta_selector.support_)      # check selected features  
  
print("Ranking: ",boruta_selector.ranking_)                # check ranking of features  
  
print("No. of significant features: ", boruta_selector.n_features_)
```

Selected Features: [True True True True True True False False True True True T
rue
True True True True False True True True True True]
Ranking: [1 1 1 1 1 1 4 3 1 1 1 1 1 1 1 2 1 1 1 1 1]
No. of significant features: 19

In [113]:

```
selected_rfe_features = pd.DataFrame({'Feature':list(X_train.columns),  
                                     'Ranking':boruta_selector.ranking_})  
selected_rfe_features.sort_values(by='Ranking')
```

Out[113]:

Feature Ranking		
0	age	1
19	ALT	1
18	AST	1
17	serum creatinine	1
15	hemoglobin	1
14	LDL	1
13	HDL	1
12	triglyceride	1
11	Cholesterol	1
20	Gtp	1
10	fasting blood sugar	1
8	systolic	1
5	eyesight(right)	1
4	eyesight(left)	1
3	waist(cm)	1
2	weight(kg)	1
1	height(cm)	1
9	relaxation	1

21	dental caries	1
16	Urine protein	2
7	hearing(right)	3
6	hearing(left)	4

In [114]:

```
X_important_train = boruta_selector.transform(np.array(X_train))
X_important_test = boruta_selector.transform(np.array(X_test))
```

In [115]:

```
# Create a new random forest classifier for the most important features
rf_important = RandomForestClassifier(random_state=1, n_estimators=1000, n_jobs = -1)

# Train the new classifier on the new dataset containing the most important features
rf_important.fit(X_important_train, Y_train)
```

Out[115]:

```
RandomForestClassifier(n_estimators=1000, n_jobs=-1, random_state=1)
```

In [117]:

```
y_important_pred = rf_important.predict(X_important_test)
rf_imp_fscore = f1_score(Y_test, y_important_pred)
```

In [118]:

```
print(rf_imp_fscore)
```

```
0.7165043432482232
```

Hyper Parameter Tunning

In [119]:

```
from sklearn.model_selection import GridSearchCV
# Create the parameter grid based on the results of random search
param_grid = {
    'bootstrap': [True, False],
    'max_depth': [5, 10, 15],
    'n_estimators': [500, 1000]}
```

In [120]:

```
rf = RandomForestClassifier(random_state = 1)

# Grid search cv
grid_search = GridSearchCV(estimator = rf, param_grid = param_grid,
                           cv = 2, n_jobs = -1, verbose = 2)
```

In [121]:

```
grid_search.fit(X_important_train, Y_train)
```

Fitting 2 folds for each of 12 candidates, totalling 24 fits

Out[121]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(random_state=1), n_jobs=-1,
             param_grid={'bootstrap': [True, False], 'max_depth': [5, 10, 15],
                          'n_estimators': [500, 1000]},
             verbose=2)
```

In [122]:

```
grid_search.best_params_
```

Out[122]:

Out[122]:

```
{'bootstrap': True, 'max_depth': 15, 'n_estimators': 1000}
```

In [124]:

```
pred = grid_search.predict(X_important_test)

f1_score(Y_test, pred)
```

Out[124]:

0.7094635848233756

In [125]:

```
imp_test_features = boruta_selector.transform(np.array(new_test_data))
```

In [126]:

```
prediction = grid_search.predict(imp_test_features)
```

In [128]:

```
res = pd.DataFrame(prediction)
res.index = new_test_data.index
res.columns = ["prediction"]

res.to_csv("prediction_results_smoke_data.csv")
files.download("prediction_results_smoke_data.csv")
```